

# LSF Job Manage System HOWTO

The High Performance Computing Center(HPCC) at SJTU  
<http://hpt.sjtu.edu.cn>

Updated at 2013-10-08

|     |  |   |
|-----|--|---|
| 1   | Check the Cluster' s Running Status via LSF            | 2 |
| 1.1 | Check the LSF Computing Nodes List <code>bhosts</code> | 2 |
| 1.2 | Check the LSF Queues <code>bqueues</code>              | 2 |
| 1.3 | Check the load of computing nodes <code>lsload</code>  | 2 |
| 2   | Submit the Jobs via LSF <code>bsub</code>              | 3 |
| 2.1 | Submit Jobs Manually                                   | 3 |
| 2.2 | Interactive Batch Submit                               | 4 |
| 2.3 | Write a LSF scripts to submit jobs                     | 4 |
| 2.4 | A more complicated LSF script                          | 5 |
| 3   | Other Job Manage Operations                            | 5 |
| 3.1 | Check the jobs' status <code>bjobs</code>              | 5 |
| 3.2 | Kill the Jobs <code>bkill</code>                       | 6 |
| 3.3 | Monitor the Output of Jobs <code>bpeek</code>          | 7 |
| 3.4 | Jobs' History Information <code>bhist</code>           | 7 |

This document will instruct you to submit and manage jobs via LSF, it contains the information about using LSF to submit, check and delete jobs.

Following the document' s operation and feedback' s methods will help you complete the job successfully. It will be very kind to provide advices, thank you!

# 1 Check the Cluster' s Running Status via LSF

## 1.1 Check the LSF Computing Nodes List bhosts

| # bhosts  |        |      |     |       |     |       |       |     |   |
|-----------|--------|------|-----|-------|-----|-------|-------|-----|---|
| HOST_NAME | STATUS | JL/U | MAX | NJOBS | RUN | SSUSP | USUSP | RSV |   |
| fat01     | ok     | -    | 16  | 0     | 0   | 0     | 0     | 0   | 0 |
| fat02     | ok     | -    | 16  | 0     | 0   | 0     | 0     | 0   | 0 |
| fat03     | ok     | -    | 16  | 0     | 0   | 0     | 0     | 0   | 0 |
| fat04     | ok     | -    | 16  | 0     | 0   | 0     | 0     | 0   | 0 |
| fat05     | ok     | -    | 16  | 0     | 0   | 0     | 0     | 0   | 0 |
| fat06     | ok     | -    | 16  | 0     | 0   | 0     | 0     | 0   | 0 |
| fat07     | ok     | -    | 16  | 0     | 0   | 0     | 0     | 0   | 0 |
| fat08     | ok     | -    | 16  | 0     | 0   | 0     | 0     | 0   | 0 |
| fat09     | ok     | -    | 16  | 0     | 0   | 0     | 0     | 0   | 0 |
| fat10     | ok     | -    | 16  | 0     | 0   | 0     | 0     | 0   | 0 |
| .....     |        |      |     |       |     |       |       |     |   |

## 1.2 Check the LSF Queues bqueues

Check the whole queues' overall information:

| # bqueues  |      |             |     |      |      |      |       |      |      |      |
|------------|------|-------------|-----|------|------|------|-------|------|------|------|
| QUEUE_NAME | PRIQ | STATUS      | MAX | JL/U | JL/P | JL/H | NJOBS | PEND | RUN  | SUSP |
| cpu        | 40   | Open:Active | -   | -    | -    | -    | 2072  | 0    | 2072 | 0    |
| fat        | 40   | Open:Active | -   | -    | -    | -    | 0     | 0    | 0    | 0    |
| gpu        | 40   | Open:Active | -   | -    | -    | -    | 288   | 0    | 288  | 0    |
| mic        | 40   | Open:Active | -   | -    | -    | -    | 0     | 0    | 0    | 0    |
| cpu-fat    | 40   | Open:Active | -   | -    | -    | -    | 16    | 0    | 16   | 0    |

Check for some queue' s information:

| # bqueues fat |      |             |     |      |      |      |       |      |     |      |
|---------------|------|-------------|-----|------|------|------|-------|------|-----|------|
| QUEUE_NAME    | PRIQ | STATUS      | MAX | JL/U | JL/P | JL/H | NJOBS | PEND | RUN | SUSP |
| fat           | 40   | Open:Active | -   | -    | -    | -    | 0     | 0    | 0   | 0    |

## 1.3 Check the load of computing nodes lsload

Check the overall load:

```
# lsload
HOST_NAME      status  r15s  r1m  r15m  ut    pg  ls    it    tmp  swp  mem
node011        ok     0.0   0.3   0.4   0%    0.0  0 49024 193G  62G  61G
node039        ok     0.0   0.6   0.5   0%    0.0  0 49024 194G  62G  61G
node041        ok     0.0   0.0   0.0   0%    0.0  0 49024 194G  62G  61G
node050        ok     0.0   0.3   0.4   0%    0.0  0 49024 194G  62G  60G
node064        ok     0.0   0.7   0.6   0%    0.0  0 49024 194G  62G  61G
node077        ok     0.0   0.7   0.5   0%    0.0  0 49024 193G  62G  61G
.....
```

Check for some node' s load:

```
# lsload node001
HOST_NAME      status  r15s  r1m  r15m  ut    pg  ls    it    tmp  swp  mem
node01         ok     0.3   0.1   0.1   1%    0.0  0  332 152G  62G  61G
```

## 2 Submit the Jobs via LSF bsub

### 2.1 Submit Jobs Manually

LSF uses bsub to submit jobs. The format of bsub is:

```
bsub -n Z -q QUEUENAME -i INPUTFILE -o OUTPUTFILE COMMAND
```

Z is the number of threads needed, -q assign the queue. If there is no option -q, the system will submit the jobs to the default queue. INPUTFILE is name of the file read by the program, OUTPUTFILE is the output file' s name.

For the serial job, COMMAND can directly be your program' s name. Example: submit the serial program mytest via LSF:

```
bsub -n 1 -q q_default -o mytest.out ./mytest
```

For the MPI parallel program, the format of COMMAND is -a mpich\_gm mpirun.lsf PROG\_NAME. Example: submit the parallel program mytest via LSF which uses 16 threads:

```
bsub -n 16 -q q_default -o mytest.out -a mpich_gm mpirun.lsf ./mytest
```

## 2.2 Interactive Batch Submit

You can start up an interactive shell environment by using `bsub` to submit multiple parallel jobs whose running arguments are the same:

```
# bsub
bsub> -n 16
bsub> -q q_default
bsub> -o output.txt
bsub> COMMAND1
bsub> COMMAND2
bsub> COMMAND3
```

It is equal to:

```
bsub -n 16 -q q_default -o output.txt COMMAND1
bsub -n 16 -q q_default -o output.txt COMMAND2
bsub -n 16 -q q_default -o output.txt COMMAND3
```

## 2.3 Write a LSF scripts to submit jobs

```
#BSUB -n 16
#BSUB -q q_default
#BSUB -o output.txt
-a mpich_gm mpirun.lsf ./mytest
```

`bsub` also accepts the state of jobs from `stdin`, that means we can write the LSF script to submit jobs. `bsub`'s script is easy to write, the code above is an example named `bsub.script`, submit `bsub.script` to LSF via input redirection:

```
bsub < bsub.script
```

It is equal to:

```
bsub -n 16 -q q_default -o output.txt -a mpich_gm mpirun.lsf ./mytest
```

## 2.4 A more complicated LSF script

```
#BSUB -J HELLO_MPI
#BSUB -o job.out
#BSUB -e job.err
#BSUB -n 256

source /lustre/utility/intel/composer_xe_2014.3.163/bin/compilervars.sh intel64
source /lustre/utility/intel/mkl/bin/intel64/mklvars_intel64.sh
source /lustre/utility/intel/impi/4.1.1.036/bin64/mpivars.sh

MPIRUN=`which mpirun`
EXE="./mpihello"

CURDIR=$PWD
cd $CURDIR
rm -f nodelist nodes >& /dev/null
touch nodelist
touch nodes
NP=0

for host in `echo $LSB_MCPU_HOSTS | sed -e 's/ /:/g' | sed 's/:n/\nn/g'`
do
echo $host >> nodelist
echo $host | cut -d ":" -f1 >> nodes
nn=`echo $host | cut -d ":" -f2`
NP=`echo $NP+$nn | bc`
done
```

## 3 Other Job Manage Operations

### 3.1 Check the jobs' status bjobs

Check the submitted jobs' running status:

```
bjobs
```

Display the jobs' running status as wide format:

```
bjobs -w
```

Display all the jobs:

```
bjobs -a
```

Display the running jobs:

```
bjobs -r
```

Display the pending jobs and reasons:

```
bjobs -p
```

Display the suspending jobs and reasons:

```
bjobs -s
```

Display detailed information of job JOBID:

```
bjobs -l JOBID
```

## 3.2 Kill the Jobs `bkill`

Kill the jobs unwanted:

```
bkill
```

Kill the job JOBID:

```
bkill JOBID
```

Remove the job JOBID from LSF instead of waiting its progresses killed by the operating system:

```
bikill JOBID
```

### 3.3 Monitor the Output of Jobs bpeek

Display the stdout and stderr output of a unfinished batch job

```
bpeek
```

Display the output of the job with the specified ID JOBID

```
bpeek JOBID
```

### 3.4 Jobs' History Information bhist

display the history of batch jobs

```
bhist
```

Display the specified job(s) JOBID only

```
bhist JOBID
```